

Fault Injection:

In this lab we will learn how to test the resiliency of an application by injecting systematic faults.

Before we start let us reset the route rules:

```
kubectl apply -f virtual-service-all-v1.yaml
```

Update the USER_NAME placeholder with your user name in `virtual-service-reviews-test-v2.yaml`:

```
vi virtual-service-reviews-test-v2.yaml
```

Then apply changes:

```
kubectl apply -f virtual-service-reviews-test-v2.yaml
```

Inject a route rule to create a fault using HTTP delay:

To start, we will inject a 7s delay between the `reviews v2` and `ratings` service for your user. `reviews v2` service has a 10s hard-coded connection timeout for its calls to the `ratings` service.

```
vi virtual-service-ratings-test-delay.yaml
```

Then apply changes:

```
kubectl apply -f virtual-service-ratings-test-delay.yaml
```

To confirm the rule is in place:

```
kubectl get virtualservice ratings -o yaml
```

Output:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
```

```

{"apiVersion":"networking.istio.io/v1alpha3","kind":"VirtualService","metadata":{"annotations":{},"name":"ratings","namespace":"default"},"spec":{"hosts":["ratings"],"http":[{"fault":{"delay":{"fixedDelay":"7s","percent":100}},"match":[{"headers":{"end-user":{"exact":"USER_NAME"}}}], "route":[{"destination":{"host":"ratings","subset":"v1"}}]}, {"route":[{"destination":{"host":"ratings","subset":"v1"}}]}]}
  creationTimestamp: 2018-10-26T15:21:42Z
  generation: 1
  name: ratings
  namespace: default
  resourceVersion: "14790"
  selfLink: /apis/networking.istio.io/v1alpha3/namespaces/default/virtualservices/ratings
  uid: d7d7347f-d932-11e8-88c5-0242f983c5dd
spec:
  hosts:
  - ratings
  http:
  - fault:
      delay:
        fixedDelay: 7s
        percent: 100
      match:
      - headers:
          end-user:
            exact: USER_NAME
      route:
      - destination:
          host: ratings
          subset: v1
  - route:
      - destination:
          host: ratings
          subset: v1

```

Now we login to /productpage as your user and observe that the page loads but because of the induced delay between services the reviews section will show:

Error fetching product reviews!

Sorry, product reviews are currently unavailable for this book.

Open the Developer Tools menu (F12) -> Network tab - webpage actually loads in about 6 seconds.

BookInfo Sample

Summary: [Wikipedia Summary](#): The Comedy of Errors is one of **William Shakespeare's** early plays. It is his shortest and one of his most farcical comedies, with a major part of the humour coming from slapstick and mistaken identity, in addition to puns and word play.

[Book Details](#) [Error fetching product reviews!](#)

Type: paperback
Pages: 200
Publisher: PublisherA
Language: English
ISBN-10: 1234567890
ISBN-13: 123-1234567890

Sorry, product reviews are currently unavailable for this book.

8 requests | 4.6 KB transferred | Finish: 6.59 s | DOMContentLoaded: 6.49 s | Load: 6.79 s

If you logout or login as a different user, the page should load normally without any errors.

The following example introduces a 5 second delay in 10% of the requests to the ratings:v1 microservice:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - fault:
      delay:
        percent: 10
        fixedDelay: 5s
    route:
    - destination:
        host: ratings
        subset: v1
```

Inject a route rule to create a fault using HTTP abort:

Let's introduce an HTTP abort to the ratings microservices for the your user.

Update the USER_NAME placeholder with your user name in virtual-service-ratings-test-abort.yaml:

```
vi virtual-service-ratings-test-abort.yaml
```

Then apply changes:

```
kubectl apply -f virtual-service-ratings-test-abort.yaml
```

To confirm the rule is in place:

```
kubectl get virtualservice ratings -o yaml
```

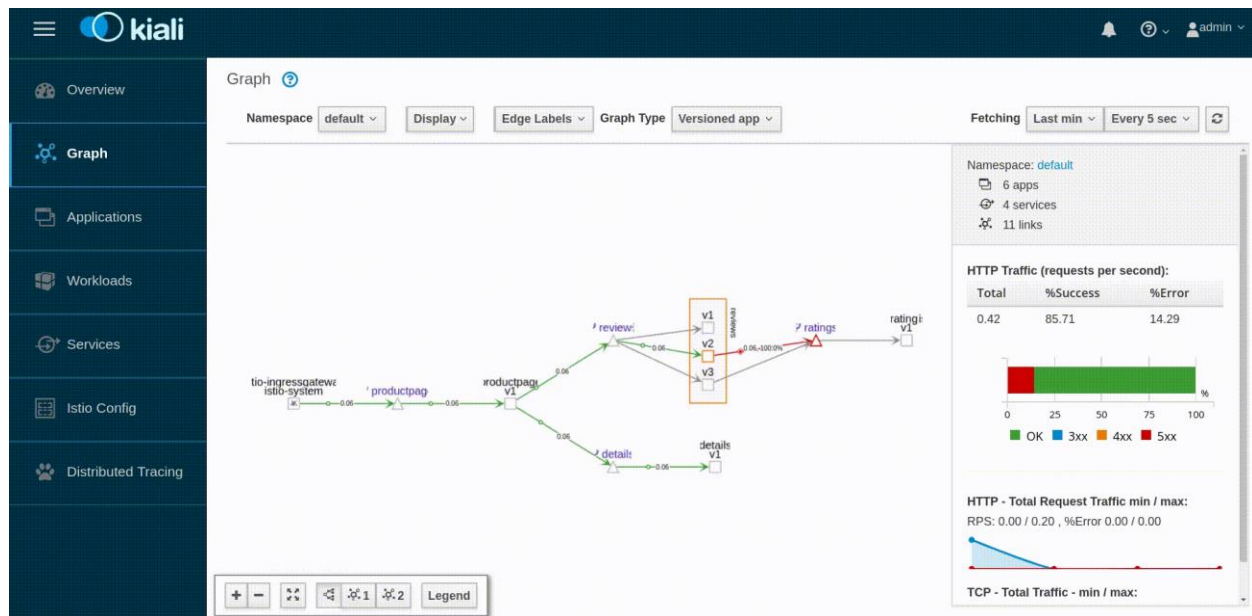
Output:

```
kind: VirtualService
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
{"apiVersion":"networking.istio.io/v1alpha3","kind":"VirtualService","metadata":{"ann
otations":{},"name":"ratings","namespace":"default"},"spec":{"hosts":["ratings"],"htt
p":[{"fault":{"abort":{"httpStatus":500,"percent":100},"match":[{"headers":{"end-
user":{"exact":"USER_NAME"}}}], "route":[{"destination":{"host":"ratings","subset":"v1
"}}]}, {"route":[{"destination":{"host":"ratings","subset":"v1"}}]}]}]}
  creationTimestamp: 2018-10-26T15:21:42Z
  generation: 1
  name: ratings
  namespace: default
  resourceVersion: "22405"
  selfLink:
/apis/networking.istio.io/v1alpha3/namespaces/default/virtualservices/ratings
  uid: d7d7347f-d932-11e8-88c5-0242f983c5dd
spec:
  hosts:
  - ratings
  http:
  - fault:
      abort:
        httpStatus: 500
        percent: 100
      match:
      - headers:
          end-user:
            exact: USER_NAME
        route:
        - destination:
            host: ratings
            subset: v1
      - route:
        - destination:
            host: ratings
            subset: v1
```

Now we login to /productpage as your user and observe that the page loads without any new delays but because of the induced fault between services the reviews section will show:

Ratings service is currently unavailable.

Check the flows in Kiali graph, where you should see the red communication between reviews:v2 and ratings.



Now, Verify the fault injection rule by logging out (or logging in as a different user), the page should load normally without any errors.

The following example returns an **HTTP 400** error code for **10%** of the requests to the ratings:v1service:

```
apiVersion: networking.istio.io/v1alpha3
kind: VirtualService
metadata:
  name: ratings
spec:
  hosts:
  - ratings
  http:
  - fault:
      abort:
        percent: 10
        httpStatus: 400
    route:
    - destination:
        host: ratings
        subset: v1
```